

**Names:**

**Raneem Ibraheem (212920896),**

**Selan Abu Saleh (212111439)**

## **Problem 1 – Template matching**

**A. First of all, we were asked to implement the `scale_down` function that takes an image and a `resize_ratio` and scales the image down by that ratio, hence we started by converting the image to the frequency representation, because when we want to scale down the image, the low frequencies that preserve the important information of the picture will be in the center, so we want to group them in the center and then start cropping the image according to the ratio and keep the middle part that we cropped. After preserving the middle part, we use the inverse transformation to go back to the spatial domain and get the scaled down image.**

```
14 def scale_down(image, resize_ratio):
15     # Your code goes here
16     # convert the image to frequency domain
17     freq_image = fftshift(fft2(image))
18     # calculate cropping indices
19     rows, cols = freq_image.shape
20     crop_start_row = int((1 - resize_ratio) * rows // 2)
21     crop_start_col = int((1 - resize_ratio) * cols // 2)
22     crop_end_row = crop_start_row + int(rows * resize_ratio)
23     crop_end_col = crop_start_col + int(cols * resize_ratio)
24     # crop and scale the frequency coefficients
25     cropped_freq = freq_image[crop_start_row:crop_end_row, crop_start_col:crop_end_col] * (resize_ratio ** 2)
26
27     # transform back to the spatial domain
28     return np.abs(fft2(fftshift(cropped_freq)))
```

**B. As for the `scale_up` image, a similar approach to the `scale_down` function is taken, but instead we create an image that is all zeros, and then we put our image that is in**

the frequency domain in the middle of the zeros image, so what happens is that we get the frequency representation of our original image padded with zeros on all the sides, of course all of these happen in relation to the ratio size. Then we do the inverse transform to go back to the spatial domain and get our scaled up image.

```
30 def scale_up(image, resize_ratio):
31     # Your code goes here
32     # convert the image to frequency domain
33     freq_image = fftshift(fft2(image))
34     # calculate dimensions for padding
35     rows, cols = freq_image.shape
36     pad_rows = int((resize_ratio - 1) * rows // 2)
37     pad_cols = int((resize_ratio - 1) * cols // 2)
38     # create a zero-padded frequency array
39     padded_freq = np.zeros((rows + 2 * pad_rows, cols + 2 * pad_cols), dtype=complex)
40     padded_freq[pad_rows:pad_rows + rows, pad_cols:pad_cols + cols] = freq_image * (resize_ratio ** 2)
41     # Transform back to the spatial domain
42     return np.abs(fft2(fftshift(padded_freq)))
43
```

C. In this section we were requested to implement the NCC function which takes the image and the pattern, and then slide the pattern across every possible position. We calculate the mean and the standard deviation for each overlapping window and normalize both to ensure the scale is similar. Then we use the formula of the NCC value at each position to get the 2d array of the NCC result where each value is the similarity between the pattern and the position in the image.

```

44 def ncc_2d(image, pattern):
45     # calculate output dimensions for the NCC result
46     img_height, img_width = image.shape
47     pat_height, pat_width = pattern.shape
48     # calculate output dimensions for the NCC result
49     output_height = img_height - pat_height + 1
50     output_width = img_width - pat_width + 1
51     # initialize the NCC result array
52     ncc_result = np.zeros((output_height, output_width))
53     # precompute the mean and normalized pattern
54     pattern_mean = np.mean(pattern)
55     normalized_pattern = pattern - pattern_mean
56     pattern_norm = np.sqrt(np.sum(normalized_pattern ** 2))
57     # compute NCC for each window in the image
58     for i in range(output_height):
59         for j in range(output_width):
60             # extract the current window from the image
61             window = image[i:i + pat_height, j:j + pat_width]
62             # compute mean and normalized window
63             window_mean = np.mean(window)
64             normalized_window = window - window_mean
65             window_norm = np.sqrt(np.sum(normalized_window ** 2))
66             # calculate the NCC value
67             ncc_result[i, j] = np.sum(normalized_window * normalized_pattern) / (window_norm * pattern_norm)
68     return ncc_result
69
70
71

```

**D. Technically in here we only had to implement the Thresholding function which takes an NCC value and a threshold and filters out the values that don't exceed the threshold.**

```

72 # a function to threshold the NCC values
73 def Thresholding(NCC, threshold):
74     exceeding = NCC > threshold
75     indices = np.argwhere(exceeding)
76     return indices
77

```

**E. For this section we had to scale up the images, and scale down the filters to an appropriate size in order to find the matching pattern, and we had use a gaussian blur for both the first image and the pattern to reduce the noise and smooth the high frequencies, and this helps to get a more accurate pattern matching. As for the second image,**

**blurring wasn't necessary because the facial features were more defined and blurring in this case might damage those features and reduce the accuracy.**

**The parameters choice in this section was through a trial-and-error process to find the best befitting parameters:**

[illegible]

```
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5> cd q1
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5\q1> python matchFaces.py
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5\q1> python matchFaces.py
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5\q1> python matchFaces.py
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5\q1> python matchFaces.py
(.venv) PS C:\Users\Alpha\Desktop\uni HW5\Image Processing\HW5\q1> python matchFaces.py
```

**Note: more trials have been conducted and there could have been better parameters to choose, but the current parameters are performing well in this case.**

```
image_scaled = scale_up(image, 1.45)
image_scaled = image_scaled - cv2.GaussianBlur(image_scaled, (15, 15), 7) + 128
pattern_scaled = scale_down(pattern, 0.8)
pattern_scaled = pattern_scaled - cv2.GaussianBlur(pattern_scaled, (19, 19), 5) + 128

display([image_scaled, pattern_scaled])

ncc = ncc_2d(image_scaled, pattern_scaled)
real_matches = Thresholding(ncc, 0.48) * (1 / 1.45)
```

```
image_scaled = scale_up(image, 1.35)
pattern_scaled = scale_down(pattern, 0.3)

display(image_scaled, pattern_scaled)

ncc = ncc_2d(image_scaled, pattern_scaled)
real_matches = Thresholding(ncc, 0.45) * (1 / 1.35)
```

Results for the last section:



**Problem 2 – Multiband blending**

**Blending two images using Laplacian pyramids is a technique that combines aspects of multi-resolution analysis and spatial blending. It enables smooth transitions between two images by mixing them at different levels of detail.**

**a:**

**We have to follow these steps:**

- **Build the Gaussian pyramids for apple image and orange image**
- **Build Laplacian pyramids to both images using Gaussian pyramids**
- **$L_i = G_i - \text{expand}(G_{i+1})$**

```
laplacian_level = cv2.subtract(gaussian_pyramid[i], resized_image)
```

I use np.float 32 to convert the input image to the float32 data type to get

**b:**

**This is how we can restore an image from a Laplacian pyramid, we begin with the smallest image which is the bottom of the pyramid and then iteratively upscale it to the size of the next larger pyramid level and add back the Laplacian detail which restores finer features, this process continues until we reach the original image size.**

```
42 def restore_from_pyramid(pyramidList, resize_ratio=2):
43     # Initialize reconstruction with the smallest image in the pyramid
44     reconstructed_image = pyramidList[-1]
45
46     # Reconstruct the image by iteratively adding upscaled images to Laplacian levels
47     for level in range(len(pyramidList) - 2, -1, -1):
48         # Determine the size of the current pyramid level
49         target_height, target_width = pyramidList[level].shape[:2]
50         # Upscale the reconstructed image to match the current level's dimensions
51         upscaled_image = cv2.resize(reconstructed_image, (target_width, target_height), interpolation=cv2.INTER_LINEAR)
52         # Add the Laplacian level to the upscaled image to reconstruct details
53         reconstructed_image = cv2.add(upscaled_image, pyramidList[level])
54
55     return reconstructed_image
```

**C:**

**To blend pyramids we have to:**

**.1. Create an empty mask with the same shape as the current pyramid level**

```
mask = np.zeros_like(pyr_apple[level], dtype=np.float32)
column_width = mask.shape[1] # Get the width of the current pyramid level
```

**2. set the left half of the mask to 1**

```
mask[:, :column_width // 2 - int((levels - level) * 1.5)] = 1.0
```

**3.define the start point for the blend area**

```
blend_first = column_width // 2 - int((levels - level) * 1.5)
```

**4.gradually create the blend by filling the transition area with a gradient:**

```
for column in range(blend_first, column_width // 2 + int((levels - level) * 1.5) + 1):
    i = column - blend_first
    mask[:, column] = 0.9 - 0.9 * i / (2 * (levels - level))
```

**5. blend the two pyramids at the current level using mask:**

$$\text{Lblend} = \text{mask}_i * \text{L1}_i + (1 - \text{mask}_i) * \text{L2}_i$$

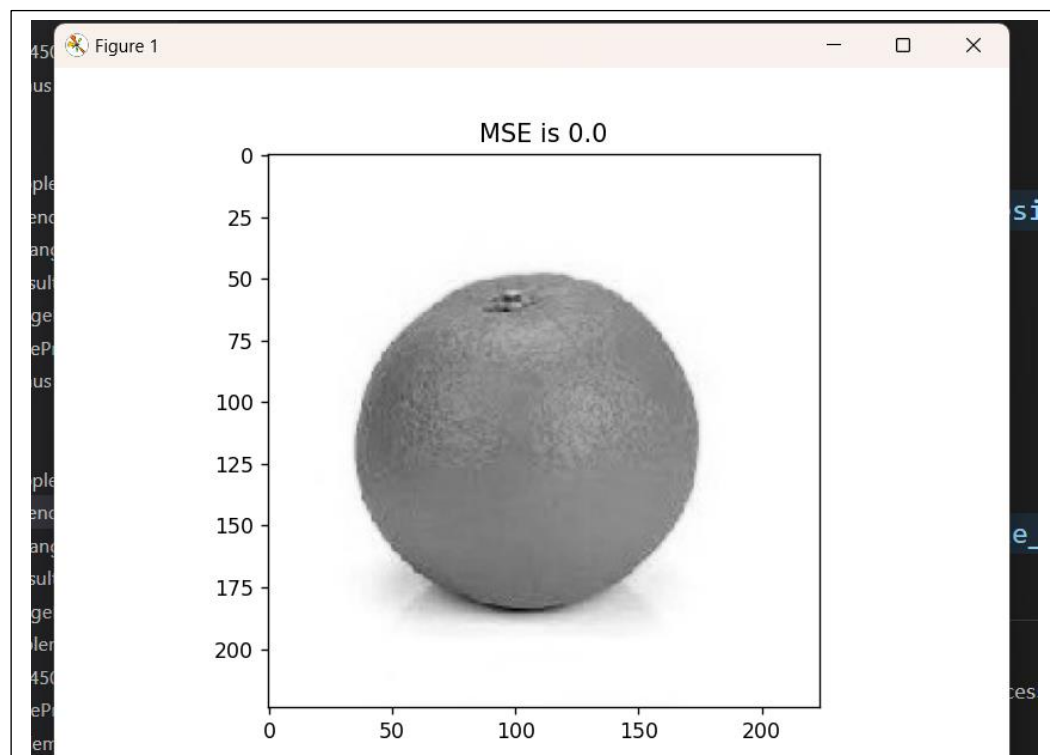
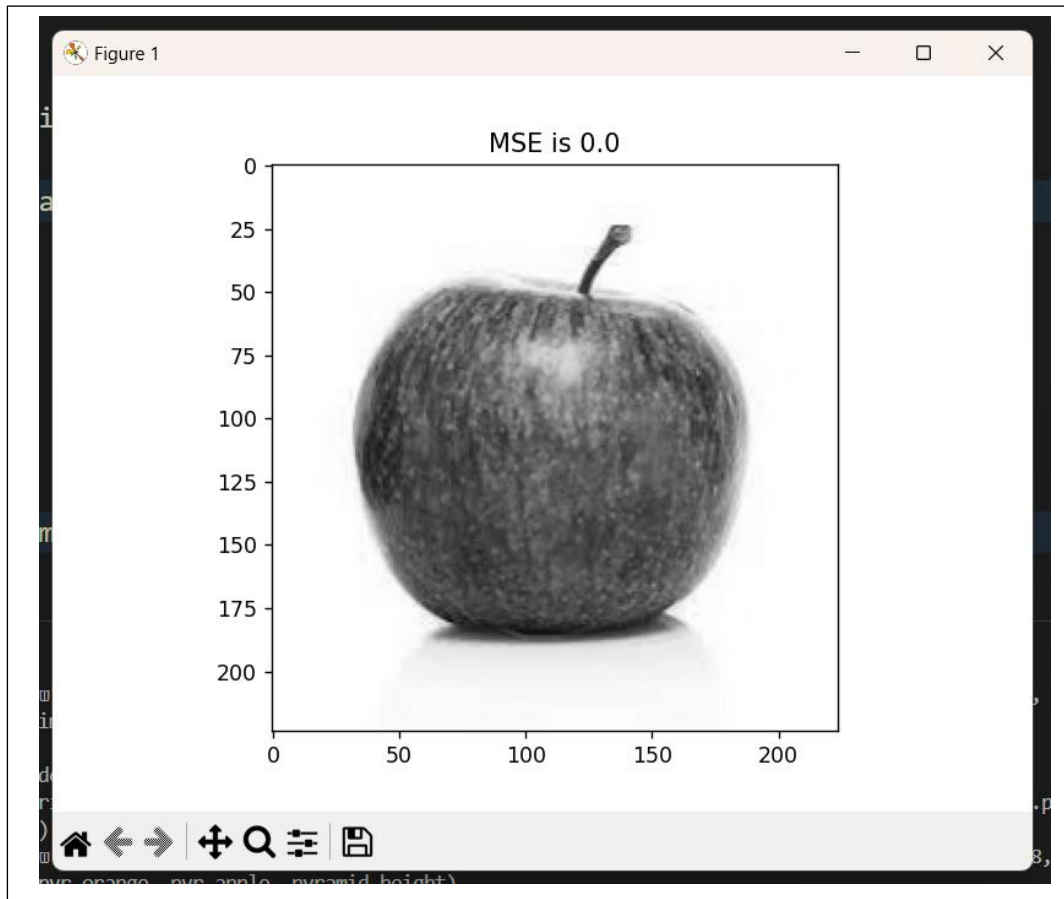
```
blended_level = pyr_orange[level] * (1 - mask) + pyr_apple[level] * mask
```

We started by defining a mask that matches the size of the current pyramid level. This mask is essentially a 2D matrix, with the same dimensions as the image at the current pyramid level, for the mask's columns, we initialized them from the first column up to  $0.5 * \text{width} - \text{current level}$ , and we set these values to 1.0, ensuring that the initial part of the mask is fully set to one image, allowing for a smooth transition to the next stage of blending.



I chose 7 levels for the pyramid to ensure more detailed and accurate blending, especially for images with a resolution of 224x224. Each level represents a progressively smaller version of the original image, capturing different levels of detail.

$224 \times 224 \rightarrow 112 \times 112 \rightarrow 56 \times 56 \rightarrow 28 \times 28 \rightarrow 14 \times 14 \rightarrow 7 \times 7 \rightarrow 3 \times 3$



**the result:**

